

MOHAMMED RIZWAN

AI/ML Engineer | Python | Machine Learning | Data Analytics
+91 9550000219 | mdrmohammedrizwan@gmail.com
[LinkedIn](#) | [GitHub](#) | [HackerRank](#)

PROFESSIONAL SUMMARY

AI/ML Engineer with hands-on experience developing Machine Learning models, Retrieval-Augmented Generation (RAG) applications, LLM-powered intelligent systems, and AI analytics solutions using Python, Scikit-learn, LangChain, FAISS, SQL and AWS. Built end-to-end AI applications involving data preprocessing, feature engineering, semantic search, vector databases, and application deployment using Streamlit. Strong foundation in Generative AI, Machine Learning, Data Analytics, and backend AI workflows with experience delivering scalable, business-oriented AI solutions.

TECHNICAL SKILLS

Programming Languages	Python, SQL
Machine Learning	Supervised Learning, Unsupervised Learning, Feature Engineering, Model Evaluation, Scikit-learn
Generative AI	LLMs, Retrieval-Augmented Generation (RAG), LangChain, LlamaIndex, Prompt Engineering, Semantic Search
Vector Databases	FAISS, Hugging Face Embeddings
Data Analytics	EDA, Data Cleaning, ETL, Data Modeling, KPI Definition
Visualization	Power BI, Tableau, Microsoft Excel
Libraries	Pandas, NumPy, Matplotlib, Seaborn
Database	MySQL
Cloud & Tools	AWS, Git, GitHub, VS Code, Jupyter Notebook, MySQL Workbench

PROFESSIONAL EXPERIENCE

- AI/ML Engineer Intern, Datavalley

Dec 2025 - Present

 - Processed 20,000+ records using Pandas and NumPy, reducing missing values by 30% before model training.
 - Built classification models using Logistic Regression and Random Forest achieving 85% validation accuracy while improving prediction consistency.
 - Developed and evaluated 15+ zero-shot, few-shot, and chain-of-thought prompt templates to improve structured information extraction from unstructured data.
 - Evaluated model performance using Accuracy, Precision, Recall, and F1-score, achieving a 15% improvement in F1-score through feature engineering and model optimization.
 - Automated Python-based ML workflows integrated with LLM APIs and automated reporting pipelines, reduced manual effort by 40%.
 - Collaborated with 4-member cross-functional teams in an Agile environment to deliver AI-driven solutions for real-world business use cases.
- Content Compliance Analyst, Wipro

Nov 2024 - Oct 2025

 - Reviewed and validated 200+ Google Ads daily while maintaining 98% policy compliance accuracy and consistently meeting SLA requirements.
 - Analyzed ad creatives and advertiser cases to identify policy violations and operational issues, using Excel reports to track performance trends and productivity metrics.

- Resolved advertiser issues through direct handling and timely escalation workflows, ensuring smooth operations and policy adherence.
- Monitored productivity metrics and adapted to evolving policy workflows, achieved 99% QA score and exceeding turnaround targets by 15%.
- Delivered data-driven insights that supported process optimization initiatives, reducing recurring operational issues by 22% and contributing to improved team efficiency.
- Consistently met 100% monthly performance KPIs across a high-volume data review environment, demonstrating sustained operational accuracy.

KEY PROJECTS

AI-Smart Attendance Tracking System

Tech Stack: Python, OpenCV, Streamlit, SQLite, Pandas, NumPy, LLM API integration (OpenAI – reporting layer), Git, GitHub

- Engineered an AI-powered attendance platform using OpenCV, SQLite, and Streamlit capable of recognizing faces with over 95% accuracy across 100+ registered users.
- Integrated LLM APIs to automatically generate attendance summaries and analytics reports, reducing manual reporting effort by 90%.
- Created interactive dashboards (Teacher, Student, Admin) supporting 100+ registered users with CSV export for reporting and analysis.
- Resolved deployment and dependency issues across multiple cloud environments, improving application stability and deployment reliability.

RAG-Powered PDF Question Answering System

Tech Stack: Python, LangChain, FAISS, HuggingFace Embeddings, Google Gemini 2.5 Flash, Streamlit, Git, GitHub

- Engineered an end-to-end RAG pipeline supporting multi-document PDF question answering through document chunking, semantic embeddings, FAISS indexing, similarity search, and Gemini-powered response generation.
- Integrated Google Gemini 2.5 Flash API with LangChain to generate accurate, context-driven responses based on retrieved document content.
- Designed an interactive Streamlit interface supporting multi-PDF uploads, chat history, source retrieval, and real-time response generation.
- Implemented prompt engineering techniques and retrieval optimization strategies to improve response relevance and reduce hallucinations.

EDUCATION

B.Tech - Computer Science and Engineering	2019 - 2023
KL University (KL Deemed to be University), Hyderabad CGPA: 8.3	
Intermediate	2017 - 2019
Narayana Junior College, Telangana Board CGPA: 8.0	

CERTIFICATIONS

- Datavalley AI/ML Engineer Internship Programme
- Aviatrix Certified Engineer - Multi-Cloud Networking Associate
- Deloitte Data Analytics Virtual Experience (Forage)
- Accenture Data Analytics and Visualization Job Simulation (Forage)
- Data Analytics - Kapil IT Skill Hub